

[Start Here](#)[Blog](#)[Books](#)[FAQ](#)[About](#)[Contact](#)

Need help with Deep Learning? [Take the FREE Mini-Course](#)

Image Augmentation for Deep Learning With Keras

by **Jason Brownlee** on June 29, 2016 in **Deep Learning**



Data preparation is required when working with neural network and deep learning models. Increasingly data augmentation is also required on more complex object recognition tasks.

In this post you will discover how to use data preparation and data augmentation with your image datasets when developing and evaluating deep learning models in Python with Keras.

After reading this post, you will know:

- About the image augmentation API provide by Keras and how to use it with your models.
- How to perform feature standardization.
- How to perform ZCA whitening of your images.
- How to augment data with random rotations, shifts and flips.
- How to save augmented image data to disk.

Let's get started.

- **Update:** The examples in this post were updated for the latest Keras API. The `datagen.next()` function was removed.
- **Update Oct/2016:** Updated examples for Keras 1.1.0, TensorFlow 0.10.0 and scikit-learn v0.18.
- **Update Jan/2017:** Updated examples for Keras 1.2.0 and TensorFlow 0.12.1.
- **Update Mar/2017:** Updated example for Keras 2.0.2, TensorFlow 1.0.1 and Theano 0.9.0.

Keras Image Augmentation API

Like the rest of Keras, the image augmentation API is simple and powerful.

Keras provides the `ImageDataGenerator` class that defines the configuration for image data preparation and augmentation. This includes capabilities such as:

- Sample-wise standardization.
- Feature-wise standardization.
- ZCA whitening.
- Random rotation, shifts, shear and flips.
- Dimension reordering.
- Save augmented images to disk.

An augmented image generator can be created as follows:

```
1 datagen = ImageDataGenerator()
```

Rather than performing the operations on your entire dataset, the generator can be iterated by the deep learning model fitting process, which reduces time. This reduces your memory overhead, but adds a small overhead.

After you have created and configured your `ImageDataGenerator`, it will calculate any statistics required to actually perform the operations. This is done by calling the `fit()` function on the data generator.

```
1 datagen.fit(train)
```

The data generator itself is in fact an iterator, returning batches of images. You can configure the batch size and prepare the data by calling the `flow()` function.

```
1 X_batch, y_batch = datagen.flow(train, train, batch_size=32)
```

Finally we can make use of the data generator. Instead of calling the `fit()` function on our model, we must call the `fit_generator()` function and pass in the data generator and the desired length of an epoch as well as the total number of epochs on which to train.

```
1 fit_generator(datagen, samples_per_epoch=len(train), epochs=100)
```

You can learn more about the Keras image data generator API in the [Keras documentation](#).

Your Start in Machine Learning

Need help with Deep Learning in Python?

Take my free 2-week email course and discover MLPs, CNNs and LSTMs (with code).

Click to sign-up now and also get a free PDF Ebook version of the course.

Start Your FREE Mini-Course Now!

Point of Comparison for Image Augmentation

Now that you know how the image augmentation API in Keras works, let's look at some examples.

We will use the MNIST handwritten digit recognition task in these examples. To begin with, let's take a look at the first 9 images in the training dataset.

```

1 # Plot images
2 from keras.datasets import mnist
3 from matplotlib import pyplot
4 # load data
5 (X_train, y_train), (X_test, y_test) = mnist.load_data()
6 # create a grid of 3x3 images
7 for i in range(0, 9):
8     pyplot.subplot(330 + 1 + i)
9     pyplot.imshow(X_train[i], cmap=pyplot.get_cmap('gray'))
10 # show the plot
11 pyplot.show()

```

Running this example provides the following image

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Example MNIST images

Feature Standardization

It is also possible to standardize pixel values across the entire dataset. This is called feature standardization and mirrors the type of standardization often performed for each column in a tabular dataset.

You can perform feature standardization by setting the `featurewise_center` and `featurewise_std_normalization` arguments on the `ImageDataGenerator` class. These are in fact set to `True` by default and creating an instance of `ImageDataGenerator` with no arguments will have the same effect.

```

1 # Standardize images across the dataset, mean=0, stdev=1
2 from keras.datasets import mnist
3 from keras.preprocessing.image import ImageDataGenerator
4 from matplotlib import pyplot
5 from keras import backend as K
6 K.set_image_dim_ordering('th')
7 # load data
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 # reshape to be [samples][pixels][width][height]
10 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
11 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
12 # convert from int to float
13 X_train = X_train.astype('float32')
14 X_test = X_test.astype('float32')
15 # define data preparation
16 datagen = ImageDataGenerator(featurewise_center=True,
17 # fit parameters from data
18 datagen.fit(X_train)
19 # configure batch size and retrieve one batch
20 for X_batch, y_batch in datagen.flow(X_train, y_train,
21 # create a grid of 3x3 images
22     for i in range(0, 9):
23         pyplot.subplot(330 + 1 + i)
24         pyplot.imshow(X_batch[i].reshape(28, 28))
25 # show the plot
26 pyplot.show()
27 break

```

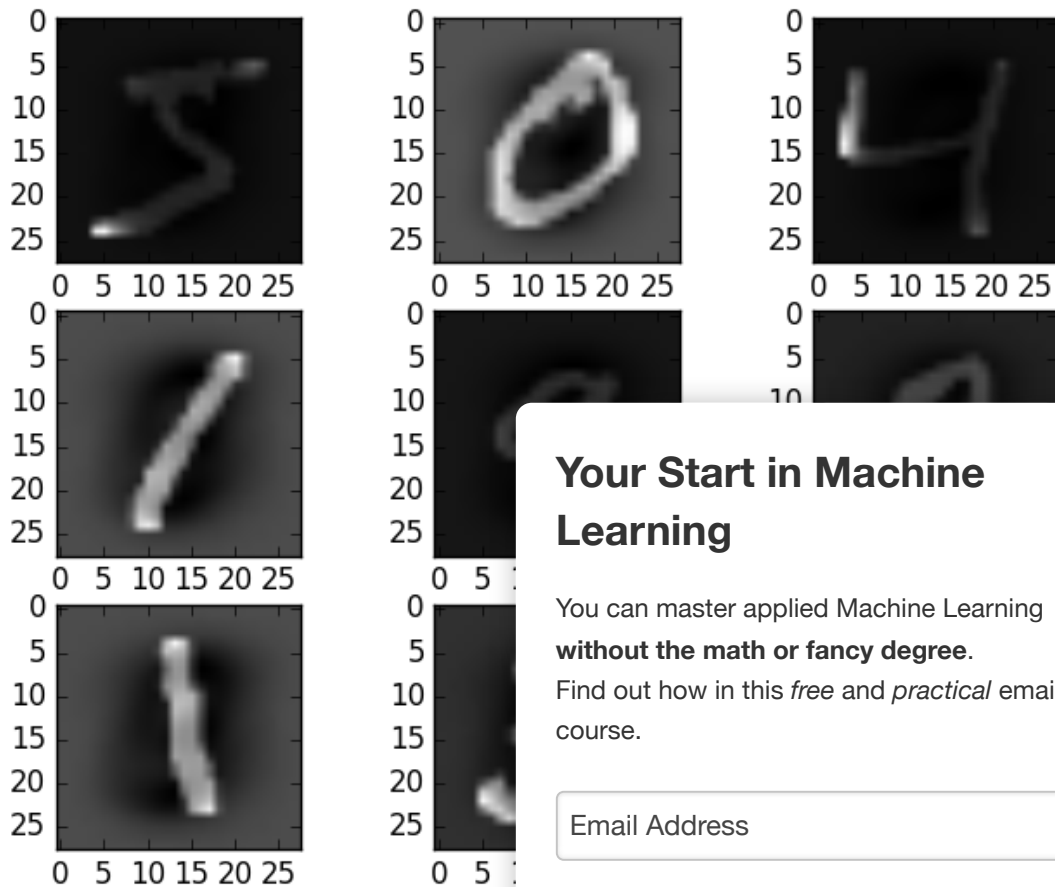
Running this example you can see that the effect is to standardize the different digits.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Your Start in Machine Learning



Standardized Fe

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

ZCA Whitening

A [whitening transform](#) of an image is a linear algebra operation that reduces the redundancy in the matrix of pixel images.

Less redundancy in the image is intended to better highlight the structures and features in the image to the learning algorithm.

Typically, image whitening is performed using the Principal Component Analysis (PCA) technique. More recently, an alternative called [ZCA](#) ([learn more in A](#)) and results in transformed images that keeps all of the original information. Transformed images still look like their originals.

You can perform a ZCA whitening transform by setting the `zca_whitening` argument to `True`.

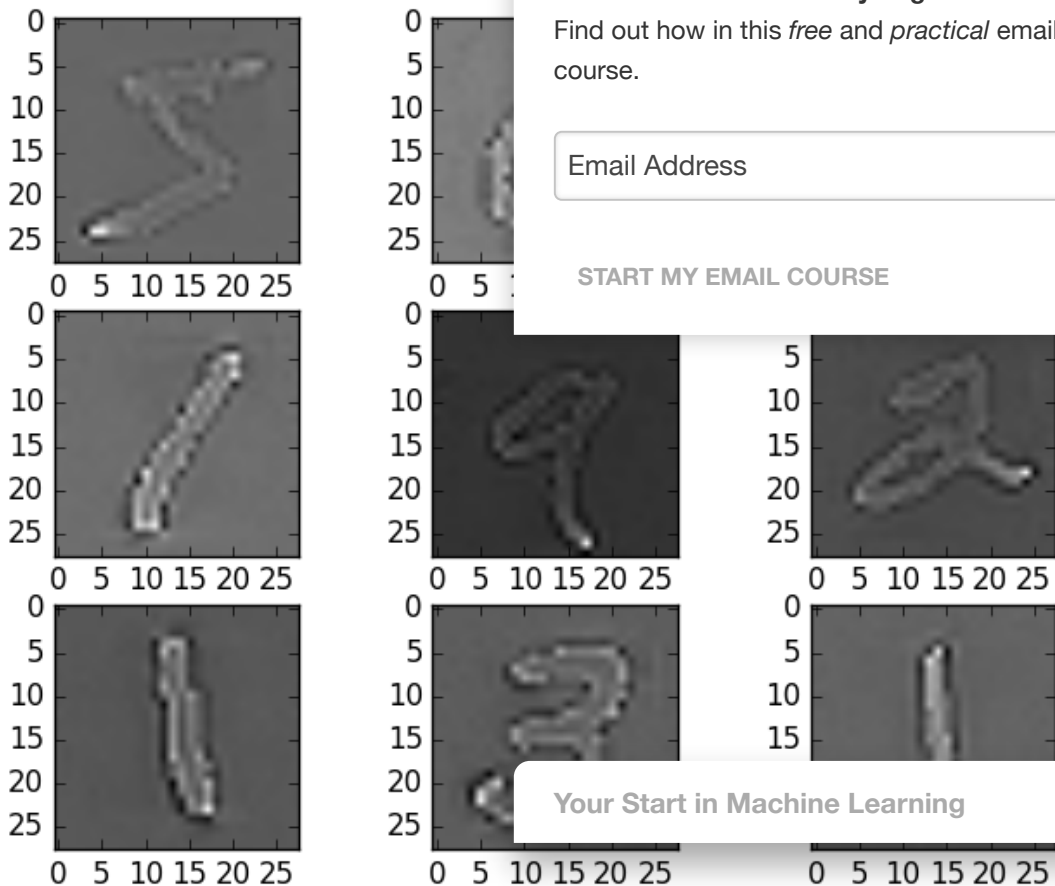
```
1 # ZCA whitening
2 from keras.datasets import mnist
3 from keras.preprocessing.image import ImageDataGenerator
4 from matplotlib import pyplot
5 from keras import backend as K
6 K.set_image_dim_ordering('th')
7 # load data
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 # reshape to be [samples][pixels][width][height]
10 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
```

```

11 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
12 # convert from int to float
13 X_train = X_train.astype('float32')
14 X_test = X_test.astype('float32')
15 # define data preparation
16 datagen = ImageDataGenerator(zca_whitening=True)
17 # fit parameters from data
18 datagen.fit(X_train)
19 # configure batch size and retrieve one batch of images
20 for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
21     # create a grid of 3x3 images
22     for i in range(0, 9):
23         pyplot.subplot(330 + 1 + i)
24         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
25     # show the plot
26     pyplot.show()
27     break

```

Running the example, you can see the same generated images as before, but each digit has been highlighted.



ZCA Whitening MNIST Images

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Your Start in Machine Learning

Random Rotations

Sometimes images in your sample data may have varying and different rotations in the scene.

You can train your model to better handle rotations of images by artificially and randomly rotating images from your dataset during training.

The example below creates random rotations of the MNIST digits up to 90 degrees by setting the `rotation_range` argument.

```

1 # Random Rotations
2 from keras.datasets import mnist
3 from keras.preprocessing.image import ImageDataGenerator
4 from matplotlib import pyplot
5 from keras import backend as K
6 K.set_image_dim_ordering('th')
7 # load data
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 # reshape to be [samples][pixels][width][height]
10 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
11 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
12 # convert from int to float
13 X_train = X_train.astype('float32')
14 X_test = X_test.astype('float32')
15 # define data preparation
16 datagen = ImageDataGenerator(rotation_range=90)
17 # fit parameters from data
18 datagen.fit(X_train)
19 # configure batch size and retrieve one batch
20 for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=10):
21     # create a grid of 3x3 images
22     for i in range(0, 9):
23         pyplot.subplot(330 + 1 + i)
24         pyplot.imshow(X_batch[i].reshape(28, 28))
25     # show the plot
26     pyplot.show()
27     break

```

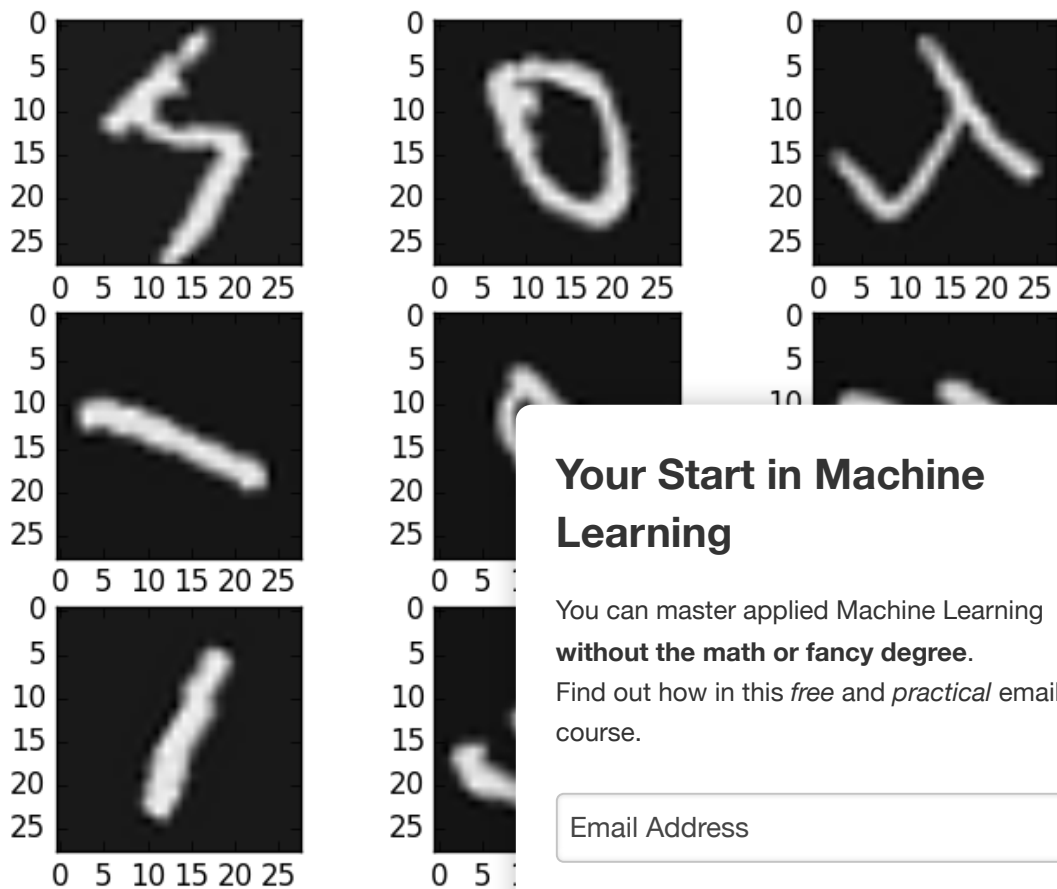
Running the example, you can see that images have been rotated up to 90 degrees. This is not helpful on this problem because the digits are not rotated, but this transform might be of help when learning to recognize digits at different orientations.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Your Start in Machine Learning



Random Rotation

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Random Shifts

Objects in your images may not be centered in the frame. They may be off-center in a variety of different ways.

You can train your deep learning network to expect and currently handle off-center objects by artificially creating shifted versions of your training data. Keras supports separate horizontal and vertical random shifting of training data by the `width_shift_range` and `height_shift_range` arguments.

```

1 # Random Shifts
2 from keras.datasets import mnist
3 from keras.preprocessing.image import ImageDataGenerator
4 from matplotlib import pyplot
5 from keras import backend as K
6 K.set_image_dim_ordering('th')
7 # load data
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 # reshape to be [samples][pixels][width][height]
10 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
11 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
12 # convert from int to float
13 X_train = X_train.astype('float32')
14 X_test = X_test.astype('float32')
15 # define data preparation
16 shift = 0.2
17 datagen = ImageDataGenerator(width_shift_range=shift, height_shift_range=shift)
18 # fit parameters from data
19 datagen.fit(X_train)

```

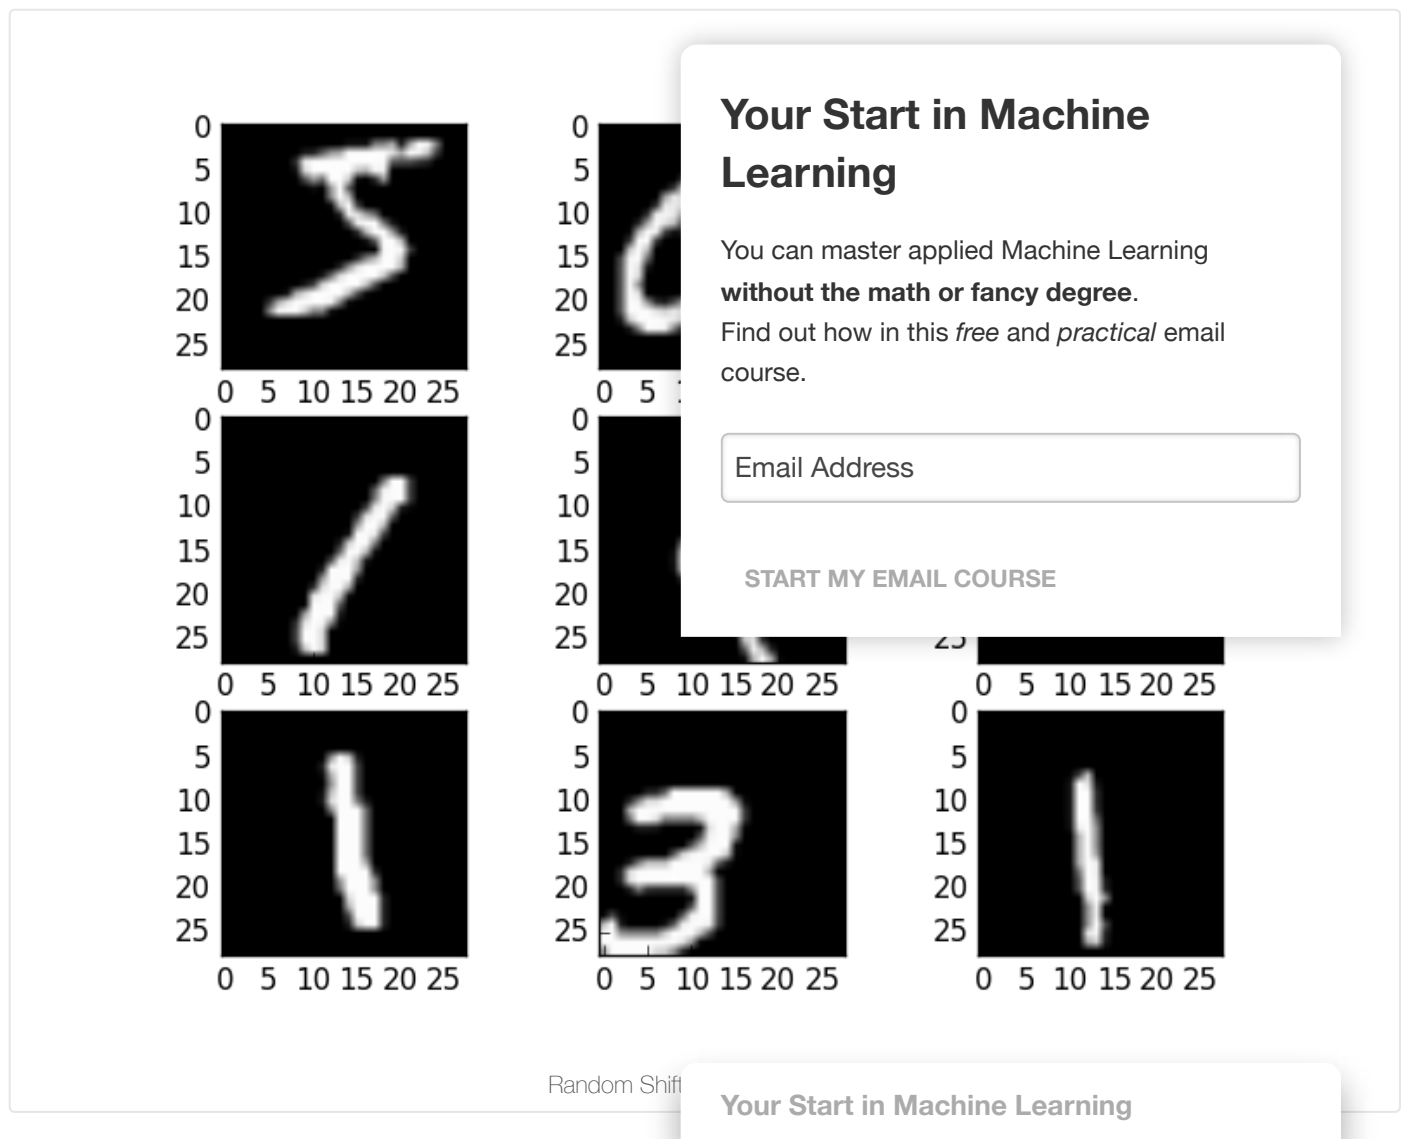
Your Start in Machine Learning


```

20 # configure batch size and retrieve one batch of images
21 for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
22     # create a grid of 3x3 images
23     for i in range(0, 9):
24         pyplot.subplot(330 + 1 + i)
25         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
26     # show the plot
27     pyplot.show()
28     break

```

Running this example creates shifted versions of the digits. Again, this is not required for MNIST as the handwritten digits are already centered, but you can see how this might be useful on more complex problem domains.



Random Flips

Another augmentation to your image data that can improve performance on large and complex problems is to create random flips of images in your training data.

Keras supports random flipping along both the vertical and horizontal axes using the `vertical_flip` and `horizontal_flip` arguments.

```

1 # Random Flips
2 from keras.datasets import mnist
3 from keras.preprocessing.image import ImageDataGenerator
4 from matplotlib import pyplot
5 from keras import backend as K

```

```

6 K.set_image_dim_ordering('th')
7 # load data
8 (X_train, y_train), (X_test, y_test) = mnist.load_data()
9 # reshape to be [samples][pixels][width][height]
10 X_train = X_train.reshape(X_train.shape[0], 1, 28, 28)
11 X_test = X_test.reshape(X_test.shape[0], 1, 28, 28)
12 # convert from int to float
13 X_train = X_train.astype('float32')
14 X_test = X_test.astype('float32')
15 # define data preparation
16 datagen = ImageDataGenerator(horizontal_flip=True, vertical_flip=True)
17 # fit parameters from data
18 datagen.fit(X_train)
19 # configure batch size and retrieve one batch of images
20 for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
21     # create a grid of 3x3 images
22     for i in range(0, 9):
23         pyplot.subplot(330 + 1 + i)
24         pyplot.imshow(X_batch[i].reshape(28, 28))
25     # show the plot
26     pyplot.show()
27     break

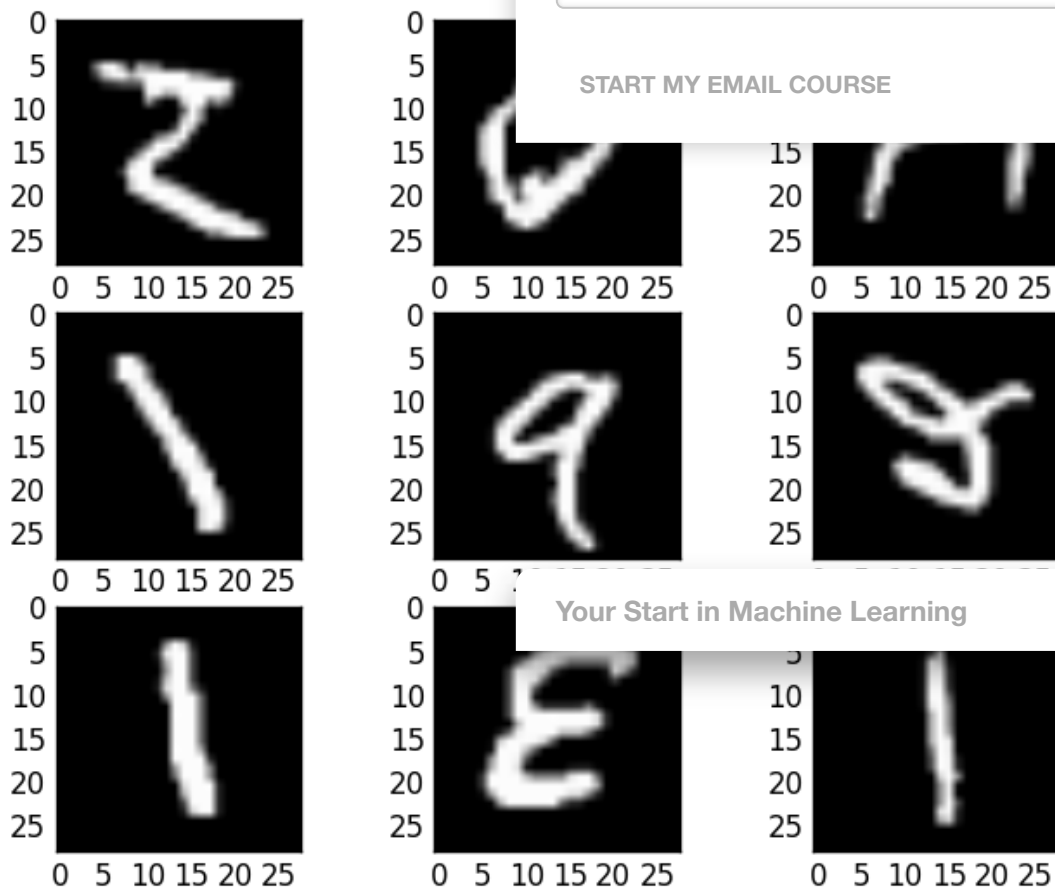
```

Running this example you can see flipped digits. For example, the digit '2' is shown in the correct left and right orientation, but this may be a scene that can have a varied orientation.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Randomly Flipped MNIST Images

Your Start in Machine Learning

Saving Augmented Images to File

The data preparation and augmentation is performed just in time by Keras.

This is efficient in terms of memory, but you may require the exact images used during training. For example, perhaps you would like to use them with a different software package later or only generate them once and use them on multiple different deep learning models or configurations.

Keras allows you to save the images generated during training. The directory, filename prefix and image file type can be specified to the `flow()` function before training. Then, during training, the generated images will be written to file.

The example below demonstrates this and writes 9 images to a “images” subdirectory with the prefix “aug” and the file type of PNG.

```

1  # Save augmented images to file
2  from keras.datasets import mnist
3  from keras.preprocessing.image import ImageDataGenerator
4  from matplotlib import pyplot
5  import os
6  from keras import backend as K
7  K.set_image_dim_ordering('th')
8  # load data
9  (X_train, y_train), (X_test, y_test) = mnist.load_data()
10 # reshape to be [samples][pixels][width][height]
11 X_train = X_train.reshape(X_train.shape[0], X_train.shape[1] * X_train.shape[2] * X_train.shape[3])
12 X_test = X_test.reshape(X_test.shape[0], X_test.shape[1] * X_test.shape[2] * X_test.shape[3])
13 # convert from int to float
14 X_train = X_train.astype('float32')
15 X_test = X_test.astype('float32')
16 # define data preparation
17 datagen = ImageDataGenerator()
18 # fit parameters from data
19 datagen.fit(X_train)
20 # configure batch size and retrieve one batch of images
21 os.makedirs('images')
22 for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9, save_to_dir='images', save_prefix='aug', save_format='png'):
23     # create a grid of 3x3 images
24     for i in range(0, 9):
25         pyplot.subplot(330 + 1 + i)
26         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
27     # show the plot
28     pyplot.show()
29     break

```

Your Start in Machine Learning

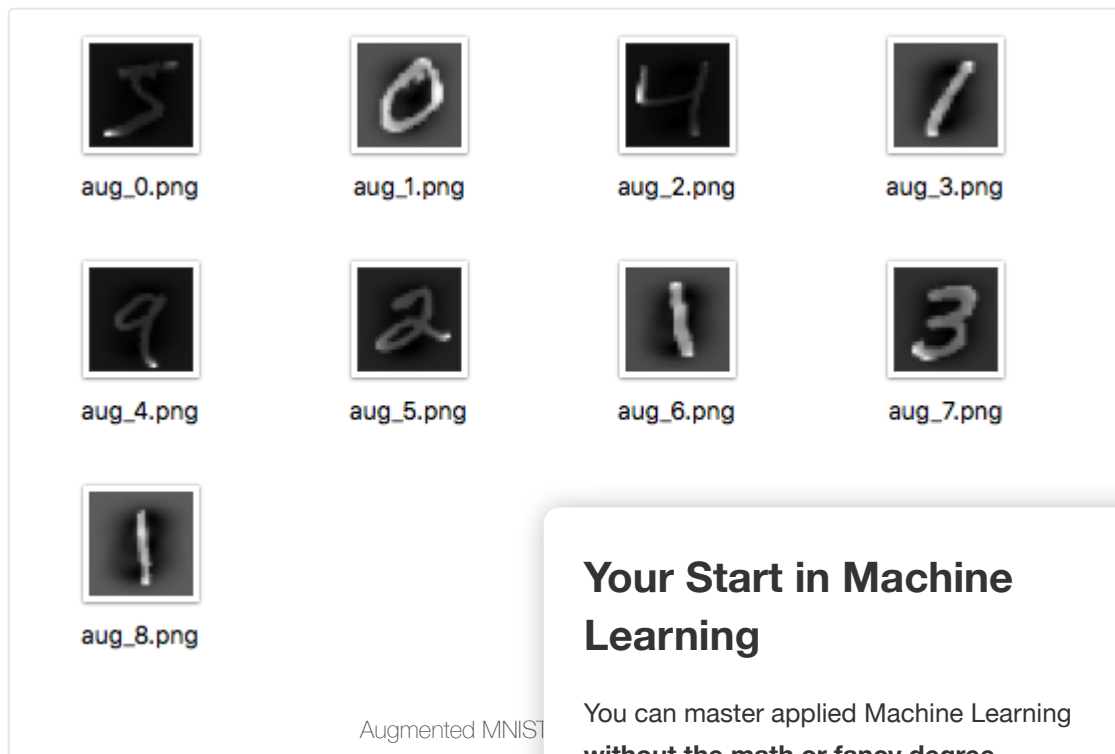
You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Running the example you can see that images are only written when they are generated.

Your Start in Machine Learning



Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Tips For Augmenting Image Data

Image data is unique in that you can review the data and get an idea of how the model may be perceive it before training.

Below are some tips for getting the most from image data learning.

- **Review Dataset.** Take some time to review your dataset in great detail. Look at the images. Take note of image preparation and augmentations that might benefit the training process of your model, such as the need to handle different shifts, rotations or flips of objects in the scene.
- **Review Augmentations.** Review sample images after the augmentation has been performed. It is one thing to intellectually know what image transforms you are using, it is a very different thing to look at examples. Review images both with individual augmentations you are using as well as the full set of augmentations you plan to use. You may see ways to simplify or further enhance your model training process.
- **Evaluate a Suite of Transforms.** Try more than one image data preparation and augmentation scheme. Often you can be surprised by results that would be beneficial.

Your Start in Machine Learning

Summary

In this post you discovered image data preparation and augmentation.


You discovered a range of techniques that you can use easily in Python with Keras for deep learning models. You learned about:

- The ImageDataGenerator API in Keras for generating transformed images just in time.
- Sample-wise and Feature wise pixel standardization.
- The ZCA whitening transform.

- Random rotations, shifts and flips of images.
- How to save transformed images to file for later reuse.

Do you have any questions about image data augmentation or this post? Ask your questions in the comments and I will do my best to answer.

Frustrated With Your Progress In Deep Learning?



What If You Could Develop A Network in Minutes

Discover
It covers self-
Multilayer Perce

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



About Jason Brownlee

Jason Brownlee, Ph.D. is a machine learning specialist who teaches developers how to get results with modern machine learning methods via hands-on tutorials.

[View all posts by Jason Brownlee →](#)

Your Start in Machine Learning

< Standard Machine Learning Datasets To Practice in Weka

How to Better Understand Your Machine Learning Data in Weka >

104 Responses to *Image Augmentation for Deep Learning With Keras*



Andy August 2, 2016 at 7:34 am #

REPLY ↩

Interesting tutorial.

I'm working through the step to standardize images across the dataset and run into the following error:

```
AttributeError Traceback (most recent call last)
in ()
18 datagen.flow(X_train, y_train, batch_size=9)
19 # retrieve one batch of images
--> 20 X_batch, y_batch = datagen.next()
21 # create a grid of 3x3 images
22 for i in range(0, 9):
```

AttributeError: 'ImageDataGenerator' object has no attribute 'next'

I have checked the Keras documentation and see

Perhaps I'm missing something.

Thanks for the great tutorials!



Jason Brownlee August 2, 2016 at 8:21 a

Yep, the API has changed. See:
<https://keras.io/preprocessing/image/>

I will update all of the examples ASAP.

UPDATE: I have updated all examples in this post to solve the problems at all.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Andy August 3, 2016 at 9:18 am #

REPLY ↩

Works like a charm! Thanks



Jason Brownlee August 3, 2016 at 9:35 am #

REPLY ↩

Glad to hear it Andy.

Your Start in Machine Learning



narayan August 9, 2016 at 6:38 pm #

REPLY ↩

for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):

File "/usr/local/lib/python2.7/dist-packages/keras/preprocessing/image.py", line 475, in next

x = self.image_data_generator.random_transform(x.astype('float32'))

File "/usr/local/lib/python2.7/dist-packages/keras/preprocessing/image.py", line 346, in

random_transform

fill_mode=self.fill_mode, cval=self.cval)

File "/usr/local/lib/python2.7/dist-packages/keras/preprocessing/image.py", line 109, in apply_transform

```
x = np.stack(channel_images, axis=0)
```

AttributeError: 'module' object has no attribute 'stack'

how to solve this error ...?



Jason Brownlee August 15, 2016 at 11:13 am #

REPLY ↩

I have not seen an error like that before. Perhaps there is a problem with your environment?

Consider re-installing Theano and/or Keras.



narayan August 26, 2016 at 9:02 pm #

i solved this error by updating theano
means it should be more than 1.9.0



Jason Brownlee August 27, 2016 at 11:13 am #

Great, glad to hear it narayan.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

REPLY ↩



narayan August 26, 2016 at 9:05 pm #

Now i have question that how to decide value for this parameter So that i can get good testing accuracy ..i have training dataset with 110 category with 32000 images ..

```
featurewise_center=False,
samplewise_center=False,
featurewise_std_normalization=False,
samplewise_std_normalization=False,
zca_whitening=False,
rotation_range=0.,
width_shift_range=0.,
height_shift_range=0.,
shear_range=0.,
zoom_range=0.,
channel_shift_range=0.,
fill_mode='nearest',
cval=0.,
horizontal_flip=False,
vertical_flip=False,
rescale=None,
dim_ordering=K.image_dim_ordering()
```

Waiting for your positive reply...

Your Start in Machine Learning



Jason Brownlee August 27, 2016 at 11:34 am #

REPLY ↩

My advice is to try a suite of different configurations and see what works best on your problem.



Walid Ahmed November 9, 2016 at 2:08 am #

REPLY ↩

Thanks a lot.

all worked fine except the last code to save images to file. I got the following exception

Walids-MacBook-Pro:DataAugmentation walidahm
Using TensorFlow backend.

Traceback (most recent call last):

File "augment_save_to_file.py", line 20, in
for X_batch, y_batch in datagen.flow(X_train, y_train,
save_prefix='aug', save_format='png'):

File "/usr/local/lib/python2.7/site-packages/keras/p
img = array_to_img(batch_x[i], self.dim_ordering, s
File "/usr/local/lib/python2.7/site-packages/keras/p
raise Exception('Unsupported channel number: ', x
Exception: ('Unsupported channel number: ', 28)

Any advice?

thanks again

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 9, 2016 at 9:52 am #

REPLY ↩

Double check your version of Keras is 1.1.0 and TensorFlow is 0.10.



Sudesh November 11, 2016 at 9:37 pm #

REPLY ↩

Hello Jason,

Thanks a lot for your tutorial. It is helping me in many ways.

I had question on mask image or target Y for training image X

Can i also transform Y along with X. Helps in the case of training for segmentation

Your Start in Machine Learning



Sudesh November 15, 2016 at 5:25 am #

REPLY ↩

I managed to do it.

datagen =

ImageDataGenerator(shear_range=0.02,dim_ordering=K._image_dim_ordering,rotation_range=5,wid


```
th_shift_range=0.05, height_shift_range=0.05, zoom_range=0.3, fill_mode='constant', cval=0)

for samples in range(0,100):
    seed = rd.randint(low=10,high=100000)
    for imgs_batch in
    datagen.flow(imgs_train,batch_size=batch_size,save_to_dir='augmented',save_prefix='aug',seed=s
eed,save_format='tif'):
    print('-')
    break
    for imgs_mask_batch in datagen.flow(imgs_mask_train, batch_size=batch_size,
save_to_dir='augmented',seed=seed, save_prefix='mask_aug',save_format='tif'):
    print('|')
    break
    print((samples+1)*batch_size)
```



Addie November 29, 2016 at 6:01 am #

This is great stuff but I wonder if you could
with three channels? I am getting some really bugg

Your Start in Machine Learning

You can master applied Machine Learning
without the math or fancy degree.

Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE



Jason Brownlee November 29, 2016 at 8

Great suggestion, thanks Addie.



Lucas December 24, 2016 at 9:02 am #

REPLY ↩

I wonder what `channel_shift_range` is about. The doc says “shift range for each
channels”, but what does this actually mean? Is it adding a random value to each channel or doing
something else?



Jason Brownlee December 26, 2016 at 7

Your Start in Machine Learning

I have not used this one yet, sorry Lucas.

You could try experimenting with it or dive into the source to see what it's all about.



Indra December 26, 2016 at 5:30 pm #

REPLY ↩

Hi,

Thanks for the post. I've one question i.e., we do feature standardization in the training set, so while
testing, we need those standardized values to apply on testing images ?



Jason Brownlee December 27, 2016 at 5:22 am #

REPLY ↩

Yes Indra, any transforms like standardization performed on the data prior to modeling will also need to be performed on new data when testing or making predictions.

In the case of standardization, we need to keep track of means and standard deviations.



Dan March 11, 2017 at 11:01 pm #

REPLY ↩

Thanks again Jason. Why do we subplot 3

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 12, 2017 at 8:24 am #

This is matplotlib syntax.

The `fig, axes = plt.subplots(3, 3)` creates a grid of 3x3 images. The number `3` places the next image (left to right, top to bottom).

I hope that helps.



Vineeth March 13, 2017 at 7:52 pm #

How do I save the augmented images into a directory with a class label prefix or even better into a subdirectory of class name?



Jason Brownlee March 14, 2017 at 8:15 am #

REPLY ↩

Great question Vineeth,

You can specify any directory and filename prefix you like in the call to `flow()`

Your Start in Machine Learning



Richa March 21, 2017 at 10:45 pm #

REPLY ↩

can we augment data of a particular class. I mean images of a class which are less, to deal with the class imbalance problem.



Jason Brownlee March 22, 2017 at 8:06 am #

REPLY ↩

Great idea.

Yes, but you may need to prepare the data for each class separately.



Lebron March 26, 2017 at 4:07 pm #

REPLY ↩

Hi Jason,

Thanks for your post!

I have a question: Does this apply to image data with RGBXYZ for each pixel?

Each of my input image is of six channels including RGB and XYZ (world coordinate), which was acquired from an organized point cloud by PCL(Point Cloud Library). I am wondering whether there is a correct way to do data augmentation for my image

I think ImageDataGenerator might be correct only for the RGB image, it means camera movement indeed well.

Thanks.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 27, 2017 at 7:52 #

Hi Lebron, I believe this specific API is designed to help you devise your own similar domain-specific transformations.



Lebron March 27, 2017 at 3:51 pm #

REPLY ↩

Thanks Jason!

To confirm, do you mean image with RGB only by “3d pixel data”? And if I have more channels, I have to do all the augmentation by myself, rather than using Keras API?



Jason Brownlee March 28, 2017 at 10:00 #

REPLY ↩

Your Start in Machine Learning

Yes, I believe that to be the case, but I could be wrong.



Brian April 16, 2017 at 9:35 am #

REPLY ↩

When I use zoom_range of 0.2 and inspect the output images, it seems to zoom h and v axes independently. However I want to have a small amount of zoom variation while preserving the aspect ratio of the images.

Also, when I specify a rotation_range, the rotated images have aliasing artefacts. Is there any way to specify rotations with antialiasing?



Jason Brownlee April 17, 2017 at 5:06 am #

REPLY ↩

I'm not sure off hand.

Do you think these concerns will affect the skill of the model?



Brian April 19, 2017 at 11:15 am #

REPLY ↩

Thanks Jason,

Aspect ratio of the image is important in a lot of images I'm not so sure about, but as they are degraded further.

I can modify my own copy of the Keras code and should be able to substitute PIL's rotation for Keras.

Keep up the good work, your writing has really helped.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee April 20, 2017 at 10:06 am #

Very nice Brian.

Let me know how you go.



Joe April 23, 2017 at 4:27 pm #

REPLY ↩

Hi Brian.

The transformations in ImageGenerator are applied using [scipy.ndimage.interpolation.affine_transform](https://docs.scipy.org/doc/scipy-0.14.0/reference/generated/scipy.ndimage.interpolation.affine_transform.html) with "order" (the order of spline used for interpolation).

Change this to one for linear interpolation or higher for higher orders.

Your Start in Machine Learning



Wuchi May 4, 2017 at 5:38 pm #

REPLY ↩

Hi Jason,

Thank you for your post! Very clear!

I am trying to use ImageDataGenerator now. But if I want to apply feature standardization to unseen data in the future, I need to save the ImageDataGenerator to disk, right? Any suggestion to do it?

Thanks a lot.



Jason Brownlee May 5, 2017 at 7:29 am #

REPLY ↩

That is correct, or you can standardize manually and just save the coefficients used.



RogerLo May 17, 2017 at 2:23 pm #

REPLY ↩

Hi Jason

I using Keras 2.x 'tf' seeting.

Why I can't using

X_batch, y_batch = datagen.flow(train, train, batch,

For example :

```
from keras.datasets import mnist
```

```
from keras.preprocessing.image import ImageDataGenerator
```

```
# load data
```

```
(X_train, y_train), (X_test, y_test) = mnist.load_data
```

```
# reshape to be [samples][pixels][width][height]
```

```
X_train = X_train.reshape(X_train.shape[0], 28, 28, 1)
```

```
X_test = X_test.reshape(X_test.shape[0], 28, 28, 1)
```

```
# convert from int to float
```

```
X_train = X_train.astype('float32')
```

```
X_test = X_test.astype('float32')
```

```
# define data preparation
```

```
datagen = ImageDataGenerator(featurewise_center=True, featurewise_std_normalization=True)
```

```
# fit parameters from data
```

```
datagen.fit(X_train)
```

```
# configure batch size and retrieve one batch of images
```

```
X_batch, y_batch = datagen.flow(X_train, y_train, batch_size=9)
```

Can you tell me why?

Thanks!

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee May 18, 2017 at 8:28 am #

What error do you get exactly?

Your Start in Machine Learning



RogerLo May 19, 2017 at 4:26 pm #

REPLY ↩

Hi, Hason

The error message is :

too many values to unpack (expected 2)



Jason Brownlee May 20, 2017 at 5:35 am #

REPLY ↩

I'm sorry I have not seen this error before, I do not have any good suggestions.



N1k31t4 November 24, 2017 at 7:53 am #

REPLY ↩

load data

```
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```

I think this should be done without the brackets

```
# load data
```

```
X_train, y_train, X_test, y_test = mnist.load_data()
```

It is returning four things, but you are only accessing

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Fahad June 16, 2017 at 12:25 am #

Hi Jason,

I have training data of the shape (2000,4,100,100) with 4 channels and dtype= uint8, stored as '.npy' file. Can you help me load this data?



Jason Brownlee June 16, 2017 at 8:03 am #

REPLY ↩

You may, try it and see.



Umberto June 23, 2017 at 7:00 pm #

REPLY ↩

Hi Jason,

Since I used the `fit_generator` method instead of `fit`, how can I correctly evaluate the model or not? Is the same for `predict_generator`? I'm a little confused.

Your Start in Machine Learning



Matthew Hancock June 25, 2017 at 1:03 am #

REPLY ↩

Hi Jason,

I have a quick question about the image augmentation. I am attempting to greatly increase the size of my training data set using data augmentation in order to increase my segmentation accuracy. Does the image generator feed multiple augmentations of the same image to the model or does it just return a single augmented version instead of the original? There seems to be no way to modify the number of augmented images the Image Data Generator actually returns.

**Matthew Hancock** June 25, 2017 at 1:06 am #

REPLY ↩

Never mind, I found my answer in the Keras documentation.

**Jason Brownlee** June 25, 2017 at 6:03 am #

REPLY ↩

Glad to hear it.

**Jason Brownlee** June 25, 2017 at 6:02 am #

Great question.

From the doc: "The data will be looped over (in <https://keras.io/preprocessing/image/>)"

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE**Bojan** March 24, 2018 at 8:48 pm #

Also, if for example I have multiple images, it create a lot of different combinations of the

- original data;
- shifted data;
- rotated data;
- noisy data;
- shifted + rotated data;
- shifted + noisy data;
- shifted + rotated + noisy data, etc.

Or will it only create one set of all the transforms created together, i.e:

- shifted + rotated + noisy data only;

If it is the latter, do you have any advice as to how to Append them maybe in a list or something?

All the best,
a very good tutorial

Your Start in Machine Learning**Jason Brownlee** March 25, 2018 at 6:28 am #

REPLY ↩

It applies all of the specified transforms in creating the augmented data.



Alice July 4, 2017 at 7:30 pm #

REPLY ↩

Hello Jason,

I made the exercise with your book which I find just great!!!

The problem is: it applies on randomly chosen images instead of doing it on the same ones from the "Point of comparison" sub-chapter. And always different samples.

How could I solve this?

I must say I don't understand how it comes the "i" applies on the `pyplot.subplot` and on the `X-batch[]`.

Thank you!!

Alice



Jason Brownlee July 6, 2017 at 10:14 am #

Think of the augmented images as random samples. You can have a new dataset with lots of variations of the original images back to the original examples.

Or perhaps I misunderstand your question?

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Nathan July 25, 2017 at 6:39 pm #

I think the problem of Alice is the same. The random modifications are never the same, which is difficult to make a comparison because they change everytime.

For example :

-the first plot gives me : 5 6 3, 0 1 9, 2 3 1

- after the ZCA whitening i have : 2 3 8, 3 2 5, 0 1 7



Jason Brownlee July 26, 2017 at 7:48 am #

REPLY ↩

Yes, by design, the augmented images are different every time it is called.

This is what we want, so the model does a better job of generalizing.

What is the problem exactly, could you help me to understand please?

Your Start in Machine Learning



Antoine Simon July 25, 2017 at 10:59 pm #

REPLY ↩

Hello Jason,

I have the same problem as Alice. I think that what she was saying was that the pictures that she plot after random modifications are never the same

It looks like the 9 pictures that are plotted are chosen randomly everytime.

It would be nice if you could answer me on this problem,

Thank you !



Jason Brownlee July 26, 2017 at 7:54 am #

REPLY ↩

Yes, this is by design. This is exactly what we want from image augmentation.



john Landler July 25, 2017 at 5:45 am #

Hi,

When I run the above script, I get this error:

Using TensorFlow backend.

C:\Users\sacheu\AppData\Local\Programs\Python\packages\keras\preprocessing\image.py:653: UserWarning: Using TensorFlow backend. Expected 3 channels or 4 channels on axis 1. However, it was passed an array following the data format convention "channels_first" (i.e. expected (height, width, channels) but got (height, width, channels)).

can you please tell me how to fix it?

i think i have the latest version of the libraries. And

Thank you.

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee July 25, 2017 at 9:49 am #

REPLY ↩

You could try changing the order of the channels in code or in the Keras configuration file.

For example, in code:

```
1 from keras import backend as K
2 K.set_image_dim_ordering('th')
```

Or if this is the cause, comment it out.

Your Start in Machine Learning



Tom November 11, 2017 at 2:10 am #

REPLY ↩

I have the same problem. Nothing works.

The message is a warning, and I still get the output images, but i.e. Feature Standardization is black and white, not gray scaled. So I suppose it is not working?

When I try to comment it out or change the order from 'th' to 'tf' – it completely brakes. The message is: ... (28 channels)

I am new and any comments are welcome.



Jason Brownlee November 11, 2017 at 9:23 am #

REPLY ↩

Perhaps double check you have the latest version of Keras installed? 2.0.8 or 2.0.9?



Tom November 13, 2017 at 9:04 pm #

(C:\ProgramData\Anaconda3\)

C:\ProgramData\Anaconda3\etc\c

(C:\ProgramData\Anaconda3) C:\U

Fetching package metadata

Solving package specifications: .

All requested packages already in

packages in environment at C:\P

#

keras 2.0.6 py36_0 conda-forge

It looks like it is up to date. But..

I used this tutorial: <https://machinelearningmastery.com/machine-learning-deep-learning-augmentation/>

```
import keras
```

```
print('keras: %s' % keras.__version__)
```

gives me:

Using Theano backend.

keras: 2.0.6

but according to tutorial it should be:

Using TensorFlow backend.

keras: 1.2.1

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee November 13, 2017 at 9:18 pm #

Your Start in Machine Learning

You can change the backend used by Keras in the `~/.keras/keras.json` configuration file.



Tom November 13, 2017 at 9:18 pm #

Ok, at the moment I can say that:

– Those warning messages are just a warning. They are in each example. I do not know how to make them disappear, but it turns out that they are harmless.

- Example Feature Standardization is black and white, not gray scaled. Maybe there is an error in script? I am to novice to spot it.
- Other examples seems to work correctly despite the warning message. (So I just saw a warning, saw different output, spotted comment with similar problem and just stopped. My bad).
- I can write set “KERAS_BACKEND=tensorflow” to change the backend (I don’t know what does it means, but never mind 😊).
- You are awesome for making those tutorials. Thank You!



Jason Brownlee November 1, 2017 at 10:04 am #

Thanks.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Xiaojie Zhang July 28, 2017 at 2:45 pm #

Hi, thanks for your share. When I try to use bigger data, I found it's very very hard to get enough about 10000 pictures and 224*224 pixels, even get When try to use fit() for zca-whitening, centering, no fit() function, I never success. Will you give some advice very much!



Jason Brownlee July 29, 2017 at 8:04 am #

REPLY ↩

Are you able to use the flow_from_directory instead of loading it all into memory?
<https://keras.io/preprocessing/image/>



Willie Maddox December 3, 2017 at 2:44 am #

REPLY ↩

I have the same exact problem with M memory. The flow_from_directory() requires your data to be in a specific directory structure.

From the Keras v2.1.2 documentation...

directory: path to the target directory. It should contain one subdirectory per class.

MS-COCO and NUS-WIDE are not structured this way natively.

Also, Xiaojie was talking specifically about the fit() function, which is called before the flow (or flow_from_directory) function.



Muneer Ahmad Dedmari August 20, 2017 at 10:02 pm #

REPLY ↩

Hi Jason,

Thanks for this nice post. I have a quick question. I have large-dataset which I am loading to model using custom data-generator. I am using it in `model.fit_generator()`. Now I want to use data-augmentation. So my question is, how/where can I use keras ImageDataGenerator? Thank you very much.



Jason Brownlee August 21, 2017 at 6:06 am #

REPLY ↩

I believe this tutorial will help:

<https://blog.keras.io/building-powerful-image-generation-models.html>

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Steve September 27, 2017 at 11:31 am #

Hi

Shall we run both `fit_transform` with origin images & then we combine them into one. On the latter one how

Thanks.

Steve



Jason Brownlee September 27, 2017 at 3:49 pm #

REPLY ↩

Just the augmented images.



Steve September 28, 2017 at 1:54 am #

REPLY ↩

Thanks Jason. Bought all of your ML books & love it I Would you write one about Transfer Learning deep enough with ImageNet and few other datasets for my own purpose ?

Thanks.

Your Start in Machine Learning



Jason Brownlee September 28, 2017 at 5:27 am #

REPLY ↩

Thanks again for your support Steve.

Yes, I have a post scheduled on re-using the VGG model. It should appear on the blog soon.



Momo October 24, 2017 at 2:13 am #

REPLY ↩

Thanks for this great post !

In the Random Shifts part, can we have control on the file names ?

Is it possible to save files as:

'aug'+original_file_name+'png' ?

Thanks.



Jason Brownlee October 24, 2017 at 5:36 am #

REPLY ↩

Yes, you can control the filenames. Pe

<https://keras.io/preprocessing/image/>

Your Start in Machine Learning

You can master applied Machine Learning

without the math or fancy degree.

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Abraham Ben November 3, 2017 at 5:12 pm #

thanks for your tutorial. When I try to use :

```
pred = model.predict_generator(data_gen.flow_from_directory(directory, target_size=(input_size, input_size)))
```

I cannot get the image filename that correspond to



Jason Brownlee November 4, 2017 at 5:27 am #

REPLY ↩

The order of predictions will match the order of files in the directory I would expect.



Abraham Ben November 4, 2017 at 3:01 pm #

REPLY ↩

I have check it, however, this is not the case since I got a pretty low accuracy compared to the val_acc. I found someone solving it by setting the batch_size to 1 when use predict_generator and ImageDataGenerator.flow_from_directory, but this i

Your Start in Machine Learning



Masun November 21, 2017 at 10:42 am #

REPLY ↩

Hi Jason,

Do you have any idea about how we can apply the same idea on signals? I mean signal augmentation for Deep learning? Thank you

Jason Brownlee November 22, 2017 at 10:47 am #

REPLY ↩



Not off hand. Consider how transforms can be used on your specific data to create new patterns. E.g. even just adding random noise is a good start.



smriti December 4, 2017 at 4:57 pm #

REPLY ↩

Once the features have been centered, using `featurewise_center=True` in `keras.preprocessing.image.ImageDataGenerator()`, How can I retrieve that statistics so that I use it to preprocess the images to be used for prediction/testing during `evaluate_generator()` etc



Jason Brownlee December 4, 2017 at 4:57 pm #

Good question, I'm not sure off the cuff. You can find more information at [https://machinelearningmastery.com/get-help-](https://machinelearningmastery.com/get-help-ask-questions-on-the-machinelearningmastery-com-forum/)

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



safae January 4, 2018 at 12:50 am #

Hi Jason,

Thank you very much for all the posts you shared,

I would like to ask if do you have an idea about imputation of missing values in time series data(such as acceleration, AC voltage, ...)

Thank you



Jason Brownlee January 4, 2018 at 8:12 am #

REPLY ↩

Not at this stage, thanks for the suggestion.



Dimitris Mallios January 30, 2018 at 3:28 am #

Your Start in Machine Learning

Hi Jason, i have a question, suppose we have semantic segmentation task and we want to rotate and flip both the image and the output image labels, how do we apply the transformations? Should i concatenate the image and the labels to a homogeneous array and then apply the appropriate transformations?



Jason Brownlee January 30, 2018 at 9:55 am #

REPLY ↩

Good question, you might want more control and apply augmentation one pair at a time.



Wafa March 10, 2018 at 10:38 am #

REPLY ↩

Hi Jason,

Thank you for your nice post! How can I use augmentation to data that I have on my disk? I see you imported mnist and I could not realize how to change this for my purpose.



Jason Brownlee March 11, 2018 at 6:18 am #

REPLY ↩

You can load the data via the augmentation API and use it to make an augmented version of your dataset. This augmentation could be used to create a grid of images.



Steven March 22, 2018 at 12:27 pm #

hi Jason,

Quick question: after the images are augmented, you need to stitch these augmented files into a single file, similar to the `mnist.load_data()` function to load them into, say a `mnist_new` variable.

```
mnist_new = []  
for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=batch_size):  
    # create a grid of 3x3 images  
    for i in range(0, 9):  
        mnist_new.append(array(X_batch, y_batch))
```

I am not sure what format should be used.

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee March 22, 2018 at 2:52 pm #

REPLY ↩

Good question, I think some experimentation would be required.

Perhaps use of numpy's `hstack` and `vstack` to create a larger array from the image arrays?

Your Start in Machine Learning



Steven March 23, 2018 at 5:05 am #

REPLY ↩

Yeah, I think need some tweaks or look into Keras' `mnist_load()` function how what the data format is when writing back. This link has the original data format:

<http://yann.lecun.com/exdb/mnist/>

Scroll down to the bottom for training and test dataset structures. They used some header information for each image. I think `hstack` and `vstack` are the way to go, but need to take care of those headers (think should be easy). I am still not sure what those "xxxx" mean. I think the files are just filled by the 28x28 small image data.



Jason Brownlee March 23, 2018 at 6:14 am #

REPLY ↩

It might be easier to write your own progressive loading function. A heck of a lot simpler actually.



Tien April 8, 2018 at 2:30 pm #

REPLY ↩

Dear Jason,

I used the code to generator featurewise standardization and it was not shown in the web.

Tien

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee April 9, 2018 at 6:04 am #

How was it different?



iles April 22, 2018 at 11:38 am #

Hello Jason

i have a dataset of images (jpg not a csv file) and each class of this images is in a folder car in folder with name cat and dog with folder name dog so how can i apply a deep learning model on this data do i need image data generator ?
thank you



Jason Brownlee April 23, 2018 at 6:11 am #

REPLY ↩

I think this tutorial will help:
<https://blog.keras.io/building-powerful-image-classification-using-cnn.html>

Your Start in Machine Learning



Kien Nguyen May 8, 2018 at 1:45 pm #

REPLY ↩

Your blog is awesome. Appreciate it.



Jason Brownlee May 8, 2018 at 2:57 pm #

REPLY ↩

Thanks.



maulik May 10, 2018 at 6:01 pm #

REPLY ↩

Hello Jason,

I am a bit confused while understanding of “with image augmentation”

Assume:-

Total Training Image:- 10 (#X_train Image)

ImageDataGenerator(rotation_range=90) #Randomly Rotation

```
model.fit_generator(datagen.flow(x_train, y_train, batch_size=1),
steps_per_epoch=len(x_train) / 1, epochs=epochs)
```

Does it mean we are training model of 10 randomly
or Are we Training model 10 original images + 10 r

Your Start in Machine Learning

You can master applied Machine Learning
without the math or fancy degree.

Find out how in this *free* and *practical* email
course.

START MY EMAIL COURSE



maulik May 10, 2018 at 8:24 pm #

I am assuming , We are training model

Q.1 Am i correct ?

Here I tried data augmentation on CIFAR10 (co
But I am not getting proper output image after

<https://pasteboard.co/HkvZLT1.png>

Here is code:-

```
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator
x1=x_train[10:11,:,:,:]
y1=y_train[10:11,]

datagen = ImageDataGenerator(rotation_range=90)

print (x1.shape) #Shape (1, 32, 32, 3)
print (y1.shape) #Shape (1, 10)
plt.imshow(x1[0]) #Image 1
plt.show()

for x_batch,y_batch in datagen.flow(x1,y1, batch_size=1):
    for i in range(0, 1):
        plt.imshow(x_batch[i].reshape(32,32,3)) #Image 2
        plt.show()
    print (y_batch)
    break
```

Jason, Could you please help ?

Thank you

Your Start in Machine Learning



Jason Brownlee May 11, 2018 at 6:34 am #

REPLY ↩

10 randomly modified versions of 10 different images.



maulik May 11, 2018 at 2:19 pm #

REPLY ↩

Got it.

I Tried to do rotation of 1 image. But when i am doing plt.show(), image is not properly displayed.

Here is image

<https://pasteboard.co/HkvZLT1.png>

Here is code:-

```
import matplotlib.pyplot as plt
from keras.preprocessing.image import ImageDataGenerator

x1=x_train[10:11,:,:,:]
y1=y_train[10:11,]

datagen = ImageDataGenerator(rotation_range=90)

print (x1.shape) #Shape (1, 32, 32, 3)
print (y1.shape) #Shape (1, 10)
plt.imshow(x1[0]) #Image 1
plt.show()

for x_batch,y_batch in datagen.flow(x1,y1, batch_size=1):
    for i in range(0, 1):
        plt.imshow(x_batch[i].reshape(32,32,3)) #Image 2
        plt.show()
    print (y_batch)
    break
```

Thank you Jason

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE



Jason Brownlee May 11, 2018

Your Start in Machine Learning

I don't have any good suggestions sorry. Perhaps try posting your code and error on StackOverflow?

Leave a Reply

Name (required)

Email (will not be published) (required)

Website

[SUBMIT COMMENT](#)

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

[START MY EMAIL COURSE](#)

Welcome to Machine Learning Mastery



Hi, I'm Jason Brownlee, Ph.D.
My goal is to make developers like You


[Read More](#)

Finally Get Started With Deep Learning

Sick of the fancy math and need for super computers?
Looking for step-by-step tutorials?
Want end-to-end projects?

[Your Start in Machine Learning](#)

Get Started With Deep Learning in Python Today!




Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.** Find out how in this *free* and *practical* email course.

Email Address

START MY EMAIL COURSE

POPULAR



Multivariate Time Series Forecasting with LSTM
AUGUST 14, 2017



How to Develop a Deep Learning Photo Caption
NOVEMBER 27, 2017



How to Use Word Embedding Layers for Deep Learning with Keras
OCTOBER 4, 2017



How to Define an Encoder-Decoder Sequence-to-Sequence Model for Neural Machine Translation in Keras
OCTOBER 26, 2017



How to Develop a Neural Machine Translation System from Scratch
JANUARY 10, 2018



Why One-Hot Encode Data in Machine Learning?
JULY 28, 2017



How to Develop an Encoder-Decoder Model with Attention for Sequence-to-Sequence Prediction in Keras
OCTOBER 17, 2017



How to Reshape Input Data for Long Short-Term Memory Networks in Keras
AUGUST 30, 2017



How to Develop an Encoder-Decoder Model for Sequence-to-Sequence Prediction in Keras
NOVEMBER 2, 2017

What is the Difference Between Test and Validation Datasets?



JULY 14, 2017

You might also like...

- [Your First Machine Learning Project in Python](#)
- [Your First Neural Network in Python](#)
- [How to Setup a Python for Machine Learning](#)
- [Your First Classifier in Weka](#)
- [Your First Model for Time Series Forecasting](#)

© 2018 Machine Learning Mastery. All Rights Reserved.

[Privacy](#) | [Disclaimer](#) | [Terms of Service](#) | [Search](#) | [Newsletter](#)

Your Start in Machine Learning

You can master applied Machine Learning **without the math or fancy degree.**

Find out how in this *free* and *practical* email course.

START MY EMAIL COURSE

Your Start in Machine Learning